

# Лекция 6

## Потоки CUDA (*CUDA Streams*)

- pinned-память - закрепленные страницы памяти;
- CUDA Streams – очереди команд;
- Мульти GPU

# Возможность одновременного копирования

Device 0: "GeForce GTX 560 Ti"

CUDA Driver Version / Runtime Version      7.5 / 7.5

CUDA Capability Major/Minor version number:    2.1

.....  
**Concurrent copy and kernel execution:      Yes with 1 copy engine(s)**  
.....



# Потоки CUDA и разрешение зависимостей при распараллеливании копирования и выполнения

(I)

*Очередь копирования    Очередь выполнения*

stream0, copy a	
stream0, copy b	
блокировка	kernel0
stream0, copy c	
stream1, copy a	
stream1, copy b	
блокировка	kernel1
stream1, copy c	

(II)

*Очередь копирования    Очередь выполнения*

stream0, copy a	
stream0, copy b	
stream1, copy a	kernel0
stream1, copy b	
stream0, copy c	kernel1
stream1, copy c	

```

#define N (1024*1024)
#define FULL_DATA_SIZE (N*20)

__global__ void kernel(int* a, int* b, int* c){
    int idx=threadIdx.x+blockIdx.x*blockDim.x;
    if(idx<N){
        int idx1=(idx+1)%256;
        int idx2=(idx+2)%256;
        float as=(a[idx]+a[idx1]+a[idx2])/3.0f;
        float bs=(b[idx]+b[idx1]+b[idx2])/3.0f;
        c[idx]=(as+bs)/2;
    }
}

int main(){
    cudaDeviceProp prop;
    int whichDevice;

    cudaGetDevice(&whichDevice);
    cudaGetDeviceProperties(&prop, whichDevice);
    if(!prop.deviceOverlap){
        printf("Device does not support overlapping\n");
        return 0;
    }

    cudaEvent_t start, stop;

```

.....

Sanders J., Kandrot E.  
*CUDA by Example, an introduction to  
 general-purpose GPU programming*,  
 Addison-Wesley, 2013.

## Выделение закрепленной (*paged-locked*) памяти

```
cudaHostAlloc( (void**)&host_a, FULL_DATA_SIZE*sizeof(int),  
              cudaHostAllocDefault);
```

```
cudaHostAlloc( (void**)&host_b, FULL_DATA_SIZE*sizeof(int),  
              cudaHostAllocDefault);
```

```
cudaHostAlloc( (void**)&host_c, FULL_DATA_SIZE*sizeof(int),  
              cudaHostAllocDefault);
```

Константа ***cudaHostAllocDefault*** означает эквивалентность функций

```
__host__ cudaError_t cudaHostAlloc ( void** pHost, size_t size, unsigned int flags )
```

и

```
__host__ cudaError_t cudaMallocHost ( void** ptr, size_t size )
```

```
.....  
cudaStream_t stream;  
cudaStreamCreate(&stream );  
  
cudaEventRecord(start,0 );  
for(int i=0; i<FULL_DATA_SIZE; i+=N){  
  
    cudaMemcpyAsync(dev_a, host_a+i, N*sizeof(int), cudaMemcpyHostToDevice, stream );  
    cudaMemcpyAsync(dev_b, host_b+i, N*sizeof(int), cudaMemcpyHostToDevice, stream );  
  
    kernel<<<N/256, 256, 0, stream>>>(dev_a, dev_b, dev_c);  
  
    cudaMemcpyAsync(host_c+i, dev_c, N*sizeof(int), cudaMemcpyDeviceToHost, stream );  
}  
  
cudaStreamSynchronize( stream );  
  
cudaEventRecord(stop,0);  
cudaEventSynchronize(stop);  
.....
```

(I)

```
.....
cudaEventRecord( start, 0 );
for(int i=0; i<FULL_DATA_SIZE; i+=N*2){
////////// П Е Р В Ы Й   П О Т О К   //////////
    cudaMemcpyAsync(dev_a0, host_a+i, N*sizeof(int), cudaMemcpyHostToDevice, stream0 );
    cudaMemcpyAsync(dev_b0, host_b+i, N*sizeof(int), cudaMemcpyHostToDevice, stream0 );

    kernel<<<N/256, 256, 0, stream0>>>(dev_a0, dev_b0, dev_c0);

    cudaMemcpyAsync(host_c+i, dev_c0, N*sizeof(int), cudaMemcpyDeviceToHost, stream0 );

////////// В Т О Р О Й   П О Т О К   //////////
    cudaMemcpyAsync(dev_a1, host_a+i+N, N*sizeof(int), cudaMemcpyHostToDevice, stream1);
    cudaMemcpyAsync(dev_b1, host_b+i+N, N*sizeof(int), cudaMemcpyHostToDevice, stream1);

    kernel<<<N/256, 256, 0, stream1>>>(dev_a1, dev_b1, dev_c1);

    cudaMemcpyAsync(host_c+i+N, dev_c1, N*sizeof(int), cudaMemcpyDeviceToHost, stream1);
}

cudaStreamSynchronize( stream0 );
cudaStreamSynchronize( stream1 );

cudaEventRecord(stop,0);
cudaEventSynchronize(stop);
.....
```

## (II)

```
.....  
cudaEventRecord(start,0)  
for(int i=0; i<FULL_DATA_SIZE; i+=N*2){  
////////// П Е Р Е М Е Ж А Е М Ы Е    П О Т О К И    //////////  
  cudaMemcpyAsync(dev_a0, host_a+i, N*sizeof(int),cudaMemcpyHostToDevice, stream0 );  
  cudaMemcpyAsync(dev_a1, host_a+i+N, N*sizeof(int),cudaMemcpyHostToDevice, stream1);  
  
  cudaMemcpyAsync(dev_b0, host_b+i, N*sizeof(int), cudaMemcpyHostToDevice, stream0 );  
  cudaMemcpyAsync(dev_b1, host_b+i+N, N*sizeof(int), cudaMemcpyHostToDevice, stream1);  
  
  kernel<<<N/256, 256, 0, stream0>>>(dev_a0, dev_b0, dev_c0);  
  kernel<<<N/256, 256, 0, stream1>>>(dev_a1, dev_b1, dev_c1);  
  
  cudaMemcpyAsync(host_c+i, dev_c0, N*sizeof(int),cudaMemcpyDeviceToHost, stream0 );  
  cudaMemcpyAsync(host_c+i+N, dev_c1, N*sizeof(int), cudaMemcpyDeviceToHost, stream1);  
}  
  
cudaStreamSynchronize( stream0 );  
cudaStreamSynchronize( stream1 );  
  
cudaEventRecord(stop,0 );  
cudaEventSynchronize(stop);  
.....
```



Device 0: "GeForce GTX 560 Ti"

CUDA Driver Version / Runtime Version 7.5 / 7.5

CUDA Capability Major/Minor version number: 2.1

Total amount of global memory: 2047 MBytes (2145927168 bytes)

( 8) Multiprocessors, ( 48) CUDA Cores/MP: 384 CUDA Cores

GPU Max Clock rate: 1645 MHz (1.64 GHz)

Memory Clock rate: 2004 Mhz

.....  
Concurrent copy and kernel execution: Yes with 1 copy engine(s)

Run time limit on kernels: Yes

Integrated GPU sharing Host Memory: No  
.....

malkov@linux-5002:~/WORKSHOP/CUDA EXERCISE/CUDA STREAMS> ./1

Elapsed time: 44.1 ms

malkov@linux-5002:~/WORKSHOP/CUDA EXERCISE/CUDA STREAMS> ./2

Elapsed time: 44.4 ms

malkov@linux-5002:~/WORKSHOP/CUDA EXERCISE/CUDA STREAMS> ./3

Elapsed time: 41.6 ms

Device 0: "GeForce GTX 650"

CUDA Driver Version / Runtime Version            6.5 / 6.5

CUDA Capability Major/Minor version number:    3.0

Total amount of global memory:                2048 MBytes (2147155968 bytes)

( 2) Multiprocessors, (192) CUDA Cores/MP:    384 CUDA Cores

GPU Clock rate:                                1110 MHz (1.11 GHz)

Memory Clock rate:                            2500 Mhz

.....  
Concurrent copy and kernel execution:        Yes with 1 copy engine(s)

Run time limit on kernels:                    Yes

Integrated GPU sharing Host Memory:        No

malkov@dew:~/WORKSHOP/PROJECTS/CUDA-EXERCISE/CUDA\_STREAMS> ./1

Elapsed time: 34.1 ms

malkov@dew:~/WORKSHOP/PROJECTS/CUDA-EXERCISE/CUDA\_STREAMS> ./2

Elapsed time: 34.1 ms

malkov@dew:~/WORKSHOP/PROJECTS/CUDA-EXERCISE/CUDA\_STREAMS> ./3

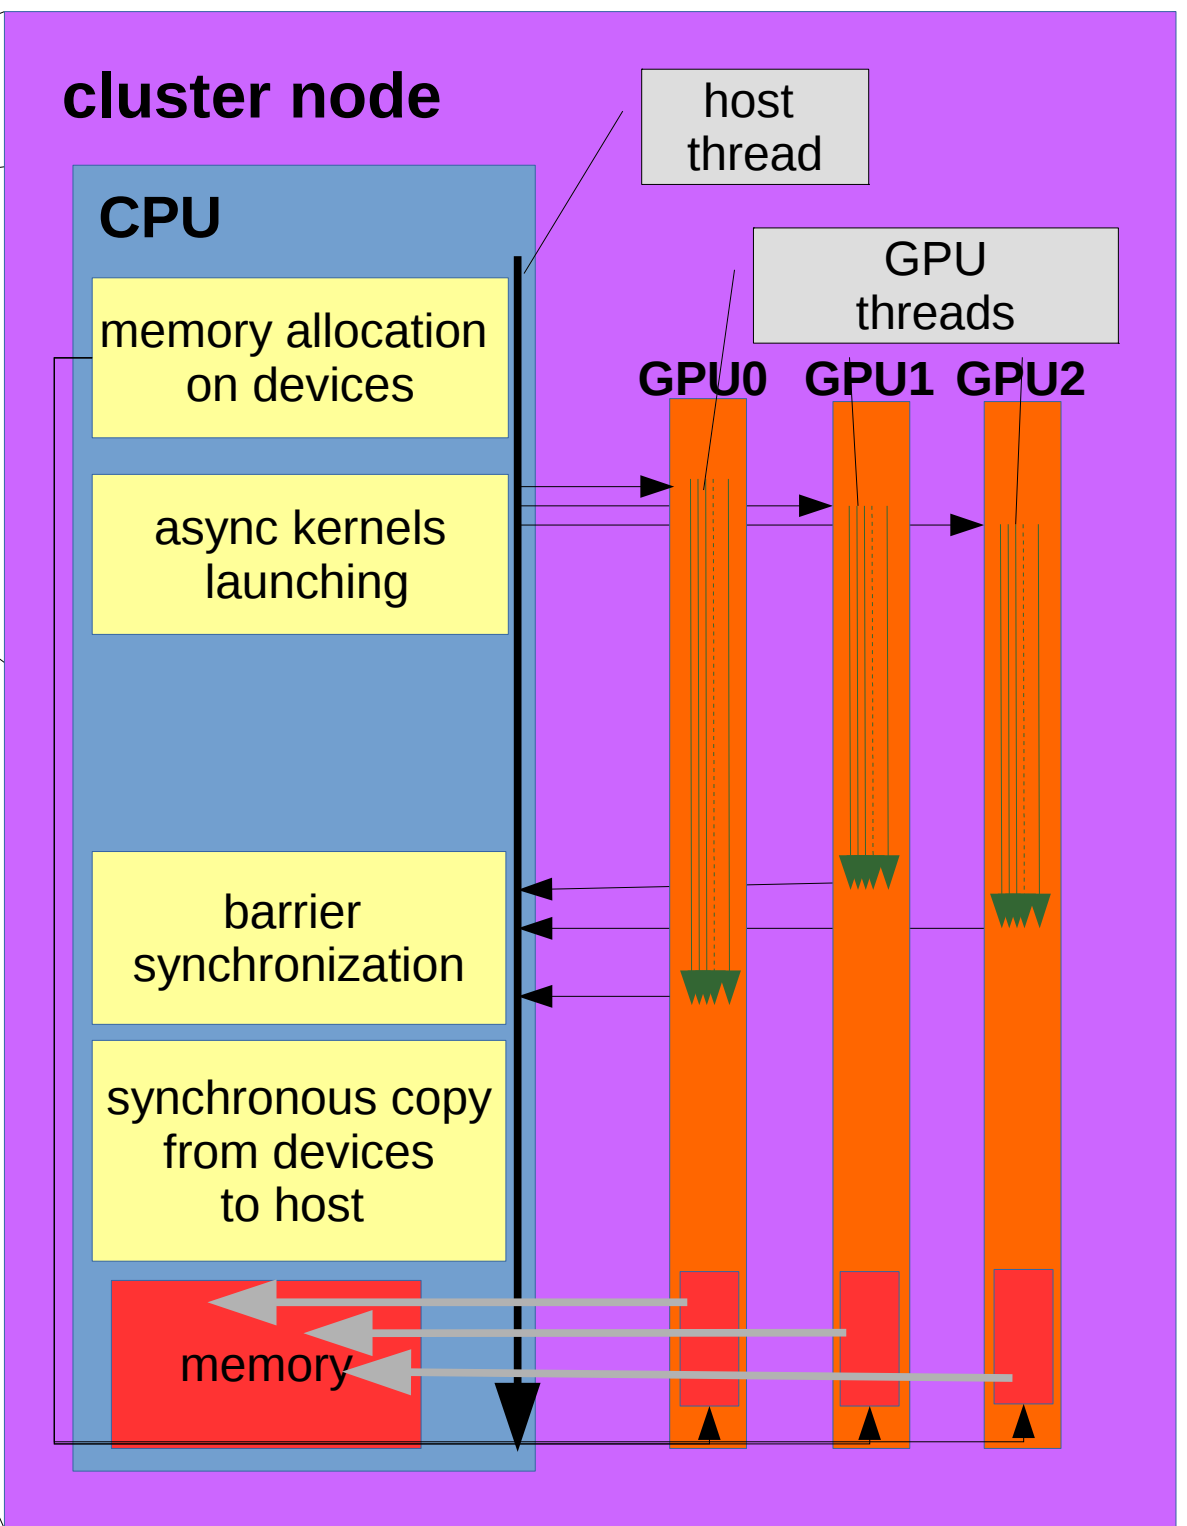

Elapsed time: 23.7 ms

malkov@dew:~/WORKSHOP/PROJECTS/CUDA-EXERCISE/CUDA\_STREAMS>

**СПАСИБО ЗА ВНИМАНИЕ**

ПРИЛОЖЕНИЕ:

select=any:ncpus=1:  
mpiprocs=1:gpus=1..3



# Выделение памяти

## ***Выделение памяти на хосте:***

```
float** Df_device=(float**)calloc(NG, sizeof(REAL*));  
float** St_device=(float**)calloc(NG, sizeof(REAL*));
```

## ***Выделение памяти на устройствах:***

```
for(int idev=0;idev< NG; idev++){  
    cudaSetDevice(assigned_devices[idev]);  
    cudaMalloc((void **) &Df_device[idev], size_of_Df/NG);  
    cudaMalloc((void **) &St_device[idev], size_of_Df/NG);  
}
```

*NG* – количество GPU.

*idev* – номер GPU.

*assigned\_devices[idev]* – идентификатор GPU.

*Df\_device[idev]*, *St\_device[idev]* – порции массивов, предназначенные для GPU с номером *idev*.

# Асинхронный запуск ядер и барьерная синхронизация.

```
cudaEvent_t *mdEventStart=(cudaEvent_t*)calloc(NG, sizeof(cudaEvent_t));  
cudaEvent_t *mdEventStop=(cudaEvent_t*)calloc(NG, sizeof(cudaEvent_t));
```

```
for(int idev=0;idev< NG; idev++){  
    cudaSetDevice(assigned_devices[idev]);  
    cudaEventCreate(&mdEventStart[idev]);  
    cudaEventCreate(&mdEventStop[idev]);  
}
```

```
for(int idev=0;idev< NG; idev++){  
    cudaSetDevice(assigned_devices[idev]);  
    cudaEventRecord( mdEventStart[idev], 0 );  
    gStossCalc<<<dim3(nvx, nvy,nvz), dim3(SPACE_CELL_OFFSET)>>>  
        (Df_device[idev], St_device[idev], N/NG, l_offset);  
    cudaEventRecord( mdEventStop[idev], 0);  
}
```

```
for(int idev=0;idev< NG; idev++){  
    cudaSetDevice(assigned_devices[idev]);  
    cudaEventSynchronize(mdEventStop[idev]);  
}
```

## Синхронное копирование с устройств на хост

```
for(int idev=0;idev< NG; idev++){  
    cudaSetDevice(assigned_devices[idev]);  
    cudaMemcpy(St+idev*size_of_DF/NG, St_device[idev],  
        size_of_Df/NG, cudaMemcpyDeviceToHost);  
}
```



# (1 GPU)

NG=1 gCalculateParams Elapsed time: 1.55024

gClearStoss gClearStoss Elapsed time: 0.213664

gTransposeDf Elapsed time: 5.2569

gStossCalc Elapsed time: **246532**

gTransposeDfInverse Elapsed time: 7.51043

gClearStoss gClearStoss Elapsed time: 0.217504

gTransposeDf Elapsed time: 5.25312

gStossCalc Elapsed time: **246536**

gTransposeDfInverse Elapsed time: 7.53414

gClearStoss gClearStoss Elapsed time: 0.217088

gTransposeDf Elapsed time: 5.25322

gStossCalc Elapsed time: **246535**

gTransposeDfInverse Elapsed time: 7.51981

gClearStoss gClearStoss Elapsed time: 0.216192

gTransposeDf Elapsed time: 5.25834

gStossCalc Elapsed time: **246533**

gTransposeDfInverse Elapsed time: 7.46163

gClearStoss gClearStoss Elapsed time: 0.216544

gTransposeDf Elapsed time: 5.25523

gStossCalc Elapsed time: **246538**

gTransposeDfInverse Elapsed time: 7.49238

gClearStoss gClearStoss Elapsed time: 0.215904

gTransposeDf Elapsed time: 5.24982

gStossCalc Elapsed time: **246536**

gTransposeDfInverse Elapsed time: 7.52288

real 24m47.123s

user 24m45.149s

sys 0m0.944s

**25,165,824**  
grid nodes  
in phase space

every loop handles  
**4,194,304** grid nodes  
in phase space

## (2 GPU)

NG=2 gCalculateParams Elapsed time: 1.53939  
NG=2 gCalculateParams Elapsed time: 1.49997  
gClearStoss gClearStoss Elapsed time: 0.211264  
gClearStoss gClearStoss Elapsed time: 0.19552  
gTransposeDf Elapsed time: 5.2217  
gTransposeDf Elapsed time: 5.23258  
gStossCalc Elapsed time: **248925**  
gStossCalc Elapsed time: **248926**  
gTransposeDfInverse Elapsed time: 7.48387  
gTransposeDfInverse Elapsed time: 7.48378  
gClearStoss gClearStoss Elapsed time: 0.200736  
gClearStoss gClearStoss Elapsed time: 0.209472  
gTransposeDf Elapsed time: 5.2145  
gTransposeDf Elapsed time: 5.22477  
gStossCalc Elapsed time: **248722**  
gStossCalc Elapsed time: **248723**  
gTransposeDfInverse Elapsed time: 7.51037  
gTransposeDfInverse Elapsed time: 7.51152  
gClearStoss gClearStoss Elapsed time: 0.199392  
gClearStoss gClearStoss Elapsed time: 0.208352  
gTransposeDf Elapsed time: 5.21789  
gTransposeDf Elapsed time: 5.22787  
gStossCalc Elapsed time: **248926**  
gStossCalc Elapsed time: **248926**  
gTransposeDfInverse Elapsed time: 7.51379  
gTransposeDfInverse Elapsed time: 7.51504  
real 12m35.874s  
user 12m33.123s  
sys 0m1.628s

**25,165,824**  
grid nodes  
in phase space

every loop handles  
**4,194,304** grid nodes  
in phase space

***two-times speedup***  
vs. 1 GPU

## (3 GPU)

NG=3 gCalculateParams Elapsed time: 1.47718  
NG=3 gCalculateParams Elapsed time: 1.53101  
NG=3 gCalculateParams Elapsed time: 1.5304  
gClearStoss gClearStoss Elapsed time: 0.140192  
gClearStoss gClearStoss Elapsed time: 0.155264  
gClearStoss gClearStoss Elapsed time: 0.124384  
gTransposeDf Elapsed time: 4.94176  
gTransposeDf Elapsed time: 5.1377  
gTransposeDf Elapsed time: 5.14787  
gStossCalc Elapsed time: **164576**  
gStossCalc Elapsed time: **164589**  
gStossCalc Elapsed time: **164590**  
gTransposeDfInverse Elapsed time: 7.12509  
gTransposeDfInverse Elapsed time: 7.43171  
gTransposeDfInverse Elapsed time: 7.44042  
gClearStoss gClearStoss Elapsed time: 0.128416  
gClearStoss gClearStoss Elapsed time: 0.13792  
gClearStoss gClearStoss Elapsed time: 0.147904  
gTransposeDf Elapsed time: 4.93661  
gTransposeDf Elapsed time: 5.13046  
gTransposeDf Elapsed time: 5.13981  
gStossCalc Elapsed time: **164575**  
gStossCalc Elapsed time: **164590**  
gStossCalc Elapsed time: **164591**  
gTransposeDfInverse Elapsed time: 7.11542  
gTransposeDfInverse Elapsed time: 7.44173  
gTransposeDfInverse Elapsed time: 7.4505  
real 5m38.815s  
user 5m36.017s  
sys 0m1.800s

**25,165,824**  
grid nodes  
in phase space

every loop handles  
**4,194,304** grid nodes  
in phase space

**superlinear speedup:**  
***aprox 5-times (not 3 !) speedup***  
vs. 1 GPU